

ECE 759
High Performance Computing for Engineering Applications
Assignment 4
Due 3/17/2025 at 23:59 PM

Submit responses to all tasks which don't specify a file name to Canvas in a file called `assignment4.pdf`. Submit all plots (if any) on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the `HW04` subdirectory on the `main` branch of your homework `git` repo with no subdirectories. Your `HW04` folder should contain `task2.cpp` and `task3.cpp`.

All commands or code must work on *Euler* without loading additional modules unless specified otherwise. A program may behave differently on your computer, so be sure to test on *Euler* before you submit. Note that this assignment is relevant to OpenMP, so the following line needs to be added to your `slurm` script:

- `#SBATCH --cpus-per-task=8` (or `-c 8` for short) should be added, which requests one node with 8 virtual cores (note the slight misnomer – Linux refers to virtual cores as cpus). The maximum number of threads required in this assignment is 8, so you should not ask for more than 8 cores.

Please submit clean code. Consider using a formatter like `clang-format`.

* Before you begin, copy the provided files from `Assignments/HW04` directory of the [ECE 759 Resource Repo](#). Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

Specify your GitHub link here:

The GitHub link to your code folder should be: <https://github.com/YourGitHubName/repo759/HW04>

In this assignment, we will simulate a dynamical system of particles interacting gravitationally, commonly referred to as the **N-body** problem. Such simulations can model systems like the orbits of planets in the Solar System or the movement of stars in a galaxy.

First, review the N-body Python code in the file located at `Assignments/HW04/nbody.py`. Below is an introduction to the N-body problem:

Force Calculation:

We will assume a system of N point particles, indexed by $i = 1, \dots, N$. Every particle has the following properties:

- mass m_i
- position $\mathbf{r}_i = [x_i, y_i, z_i]$
- velocity $\mathbf{v}_i = [vx_i, vy_i, vz_i]$

Each particle experiences the gravitational force exerted by all other particles according to Newton's law of universal gravitation, which follows the inverse-square law. The acceleration \mathbf{a}_i of particle i is given by:

$$\mathbf{a}_i = G \sum_{j \neq i} m_j \frac{\mathbf{r}_j - \mathbf{r}_i}{|\mathbf{r}_j - \mathbf{r}_i|^3}$$

where G is the Gravitational constant.

To implement this, we have a Python function that computes the gravitational force on each particle using an input matrix of size $N \times 3$ that contains the positions of all particles:

```

1 def getAcc(pos, mass, G, softening):
2     """
3     Calculate the acceleration on each particle due to Newton's Law
4     pos is an N x 3 matrix of positions
5     mass is an N x 1 vector of masses
6     G is Newton's Gravitational constant
7     softening is the softening length
8     a is N x 3 matrix of accelerations
9     """
10
11     N = pos.shape[0]
12     a = np.zeros((N, 3))
13
14     for i in range(N):
15         for j in range(N):
16             if i != j:
17                 dx = pos[j, 0] - pos[i, 0]
18                 dy = pos[j, 1] - pos[i, 1]
19                 dz = pos[j, 2] - pos[i, 2]
20                 inv_r3 = (dx**2 + dy**2 + dz**2 + softening**2) ** (-1.5)
21                 a[i, 0] += G * (dx * inv_r3) * mass[j, 0]
22                 a[i, 1] += G * (dy * inv_r3) * mass[j, 0]
23                 a[i, 2] += G * (dz * inv_r3) * mass[j, 0]
24
25     return a

```

The *softening* parameter in the code is a small value added to avoid numerical issues when two particles are very close to each other. In such cases, the acceleration predicted by the inverse-square law would approach infinity. In reality, masses are not perfect point particles and have a finite size, so this parameter prevents unrealistic behavior in the simulation by smoothing out the gravitational force at short distances.

Time Integration:

The positions and velocities of the particles are updated using a 'kick-drift-kick' method. For each timestep Δt , the following steps are applied to every particle:

1. Half-step 'kick': Update the velocity based on the current acceleration:

$$\mathbf{v}_i^{T+1} = \mathbf{v}_i^T + \frac{\Delta t}{2} \times \mathbf{a}_i^T$$

2. Full-step 'drift': Update the position using the new velocity:

$$\mathbf{r}_i^{T+1} = \mathbf{r}_i^T + \Delta t \times \mathbf{v}_i^T$$

3. Another half-step 'kick': Update the velocity again based on the new acceleration.

This evolution is implemented in the code using a for-loop and the acceleration function defined earlier:

```

1     # Simulation Main Loop
2     for i in range(Nt):
3         # (1/2) kick
4         vel += acc * dt / 2.0
5
6         # drift
7         pos += vel * dt
8
9         # ensure particles stay within the board limits
10        pos[pos > board_size] = board_size
11        pos[pos < -board_size] = -board_size
12
13        # update accelerations
14        acc = getAcc(pos, mass, G, softening)
15
16        # (1/2) kick
17        vel += acc * dt / 2.0
18
19        # update time
20        t += dt

```

Initial Conditions:

To run the simulation, we need to specify the initial positions and velocities of the particles at time $t = 0$. You can experiment with different setups in the provided code to observe how the system evolves under various initial conditions.

Problem 1. Please review the introduction to the N-body problem above. Ensure that you understand both the algorithm and the Python code.

- a) Install Python library `numpy` and `matplotlib`:
`python3 -m pip install numpy matplotlib`
- b) Run the `nbody.py` script either on your personal machine or on the Euler compute node. If you run it locally, you will see an animation of the N-body simulation. At the end of the simulation, the code will automatically save a plot. Rename this plot as `task1.png` and submit it to Canvas.

Problem 2. Port the N-body simulation code into C++. A skeleton code is provided in `Assignments/HW04/task2.cpp`.

- a) Write a program `task2.cpp` that accomplish the following:
 - Complete the provided skeleton code from `task2.cpp`.
 - Compile:
`g++ task2.cpp -Wall -O3 -std=c++17 -o task2`
 - Run the program; make sure you use `Slurm`:
`./task2 number_of_particles simulation_end_time`
 - (Optional) For debugging and visualization, we provide a function `savePositionsToCSV`. This function saves the positions at each iteration to a CSV file. You can then use `plot_positions.py` to visualize the C++ simulation results. However, if you're measuring performance, disable this function as it introduces significant I/O overhead.

Problem 3. Parallelize the N-body C++ code using OpenMP.

- a) Write a program `task3.cpp` that accomplish the following:
 - Add an argument `num_threads`, which specifies the number of threads to launch with OpenMP.
 - Insert OpenMP directives (`#pragma omp parallel for`) in sections of the code that can be parallelized.
 - Compile:
`g++ task3.cpp -Wall -O3 -std=c++17 -o task3 -fopenmp`
 - Run the program; make sure you use `Slurm` and request multiple CPU cores:
`./task3 number_of_particles simulation_end_time num_threads`

Problem 4. Let's do experiments using different OpenMP scheduling policies.

- a) In `task3.cpp`, try the following scheduling policies: `static`, `dynamic`, and `guided`.
- b) On *Euler*, via `Slurm` do the following:
For each scheduling policy, run `task3` with the following parameters:

- `number_of_particles` = 100 or larger
- `simulation_end_time` = 100 or larger
- `num_threads` = 1, 2, \dots , 8

Generate plots called `task4.pdf` that show the time taken by the N-body simulation versus `num_threads` on a linear-linear scale. The `task4.pdf` should contain three plots, one for each scheduling policy. Feel free to share the plots on Piazza.